



# Intro to (Practical) Reinforcement Learning

Dec 2019 Second Nepal Winter School in AI



# Course Format

- 4 x 45min with 3min breaks in between
- Afterwards 1h lab session with quiz + python fun
- Ask questions anytime, interrupt me!



# Goal

- Understand basic concepts around RL ( $S, A, T, R, \gamma$ )
- Understand basic algorithms (policy iteration, SARSA, etc.)
- Be able to use DRL at a grad student level\*
- NOT: understand SotA algorithms / create new SotA

\* i.e. be able to download somebody else's algorithm and run it on your task



# Outline

**Part 1 - Intro & MDPs**  
(Examples, Markov stuff)

**Part 2 - RL for Evaluation**  
(Policy Evaluation, TD(0))

**Part 3 - Model-Free RL for Control**  
(Q learning / SARSA)

**Part 4 - Practical RL**  
(OpenAI Gym, SotA algorithms, etc.)

---

# Part 1 - Introduction + MDP

# Why RL tho?



Task (e.g. is this image a cat?)

Supervised Learning

Answer:  
Yes, it's a cat

No, it's a rabbit  
(Ground Truth)

MSE/BCE


Reinforcement Learning

Answer:  
Yes, it's a cat

That answer  
was incorrect

Policy loss

# Why RL tho?



Task (e.g. make this robot stand up & walk)

Supervised Learning

Reinforcement Learning

Answer: Some movement

Lol, IDK!


Answer: Some movement

Warmer, warmer!

MSE/BCE

Policy loss

# RL Loopdy Loop








# What else can RL do?



Source: Stanford, <https://www.youtube.com/watch?v=0Jl04Jjoccc>



Source: Deepmind, <https://www.youtube.com/watch?v=TmPfTjtdgg>






# Resources


- Youtube, “RL Course by David Silver”, 11 lectures x 1h30:  
<https://www.youtube.com/playlist?list=PLzuuYNsE1EZAXYR4FJ75jcJseBmo4KQ9->
- Youtube, Abbeel & Klein, [http://ai.berkeley.edu/lecture\\_videos.html](http://ai.berkeley.edu/lecture_videos.html)
- Sutton & Barto “Reinforcement Learning: An Introduction”:  
<http://incompleteideas.net/book/bookdraft2017nov5.pdf>
- Spinning Up in DRL: <https://spinningup.openai.com/en/latest/>

Some slides/formulas were copied from the above

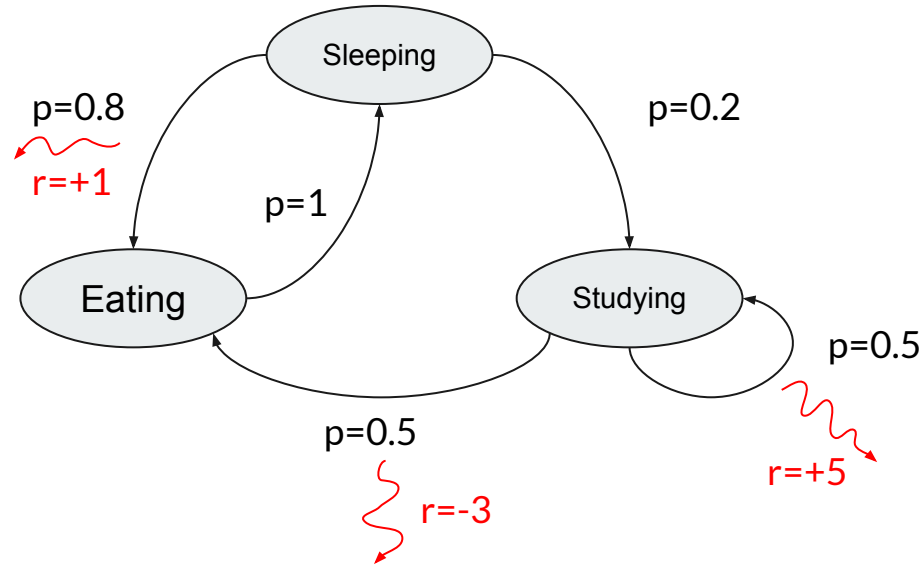
# Finite State Machine



# Markov Chain

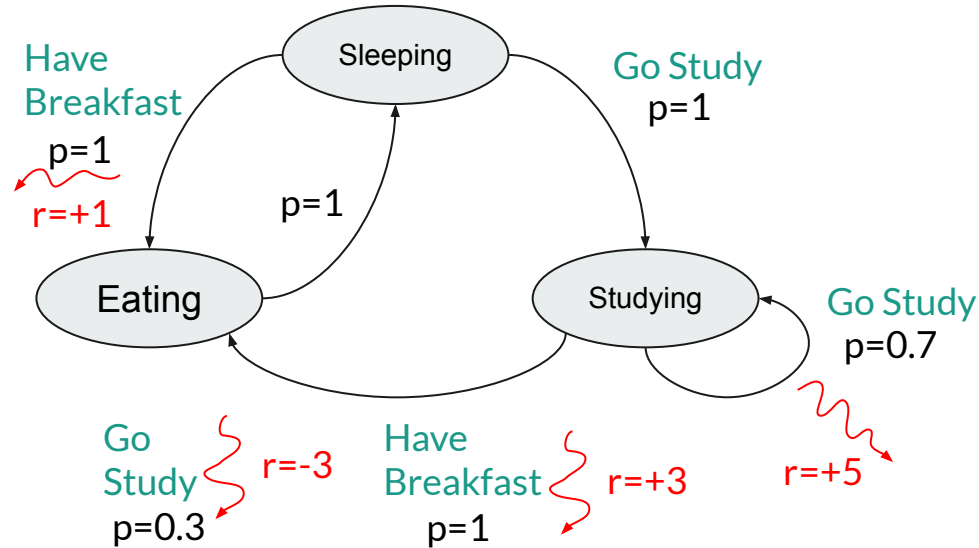


# Markov Reward Process



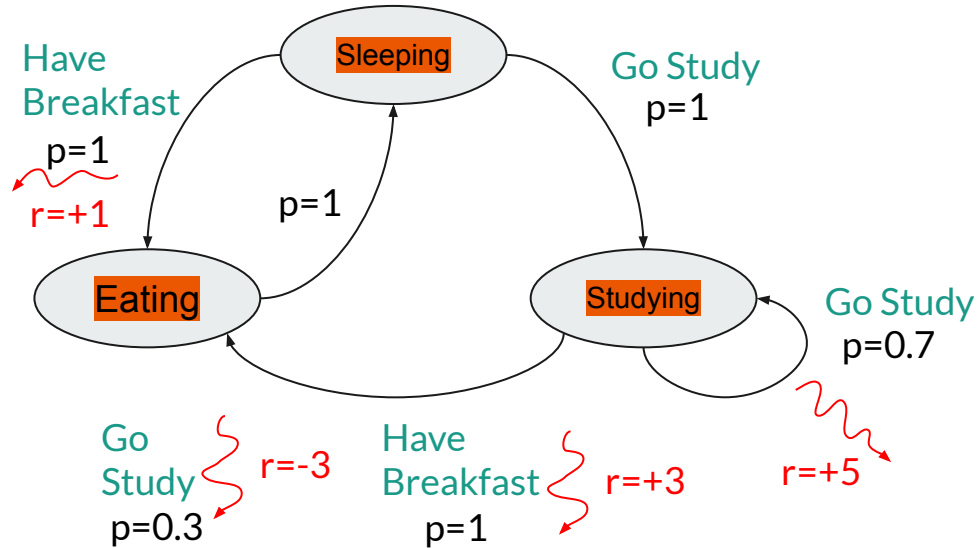


# Markov Decision Process



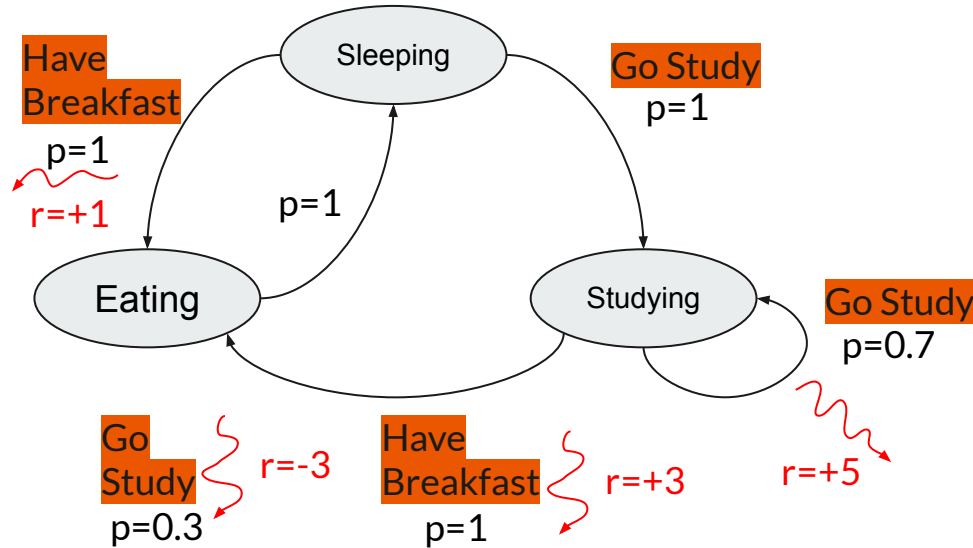
# Markov Decision Process

- States (S)



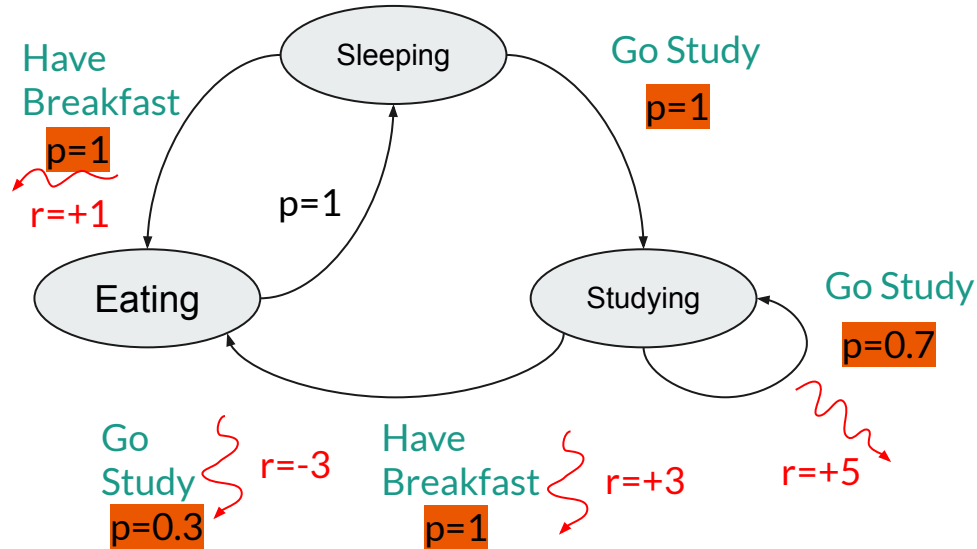
# Markov Decision Process

- States (S)
- Actions (A)



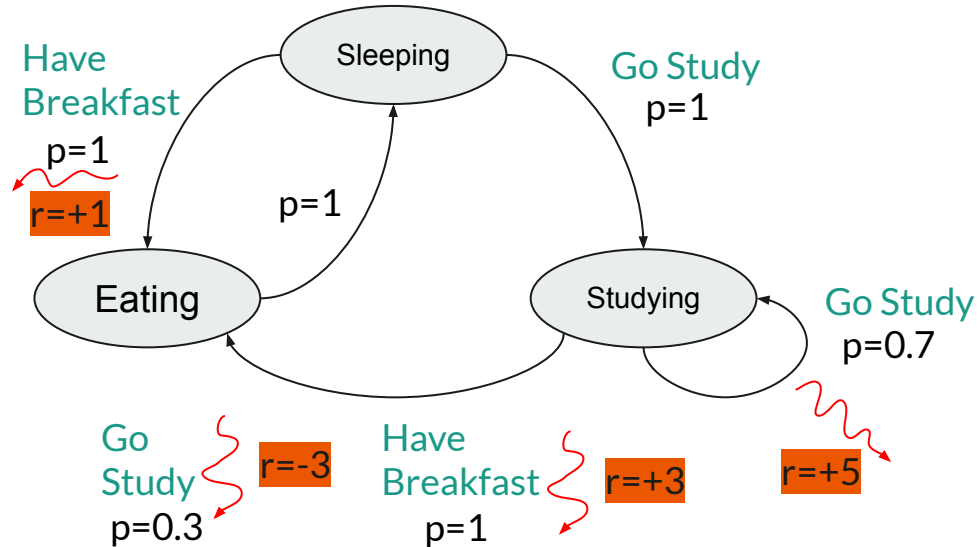
# Markov Decision Process

- States (S)
- Actions (A)
- Transition Probabilities (T/P)



# Markov Decision Process

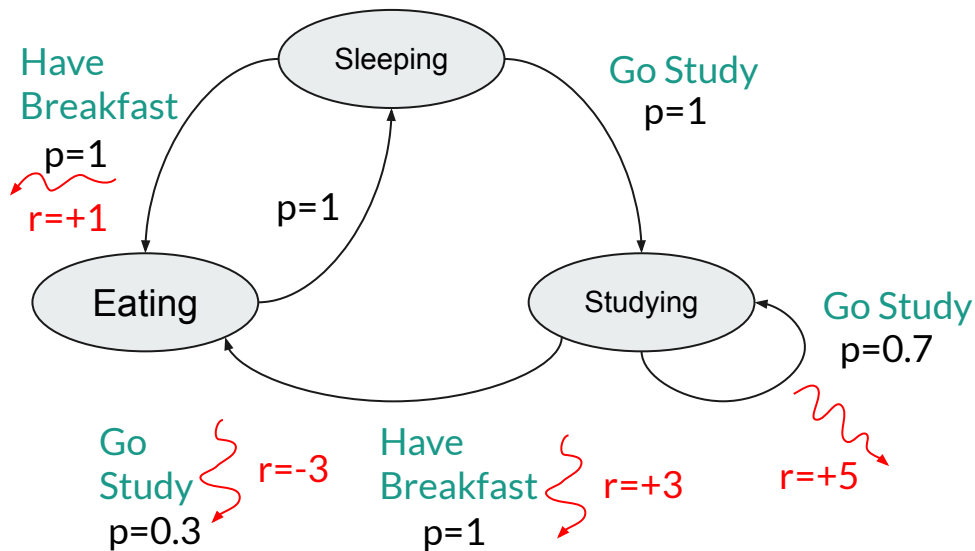
- States (S)
- Actions (A)
- Transition Probabilities (T/P)
- Rewards (R)



# Markov Decision Process

- States (S)
- Actions (A)
- Transition Probabilities (T/P)
- Rewards (R)
- Discount Factor ( $\gamma$ )

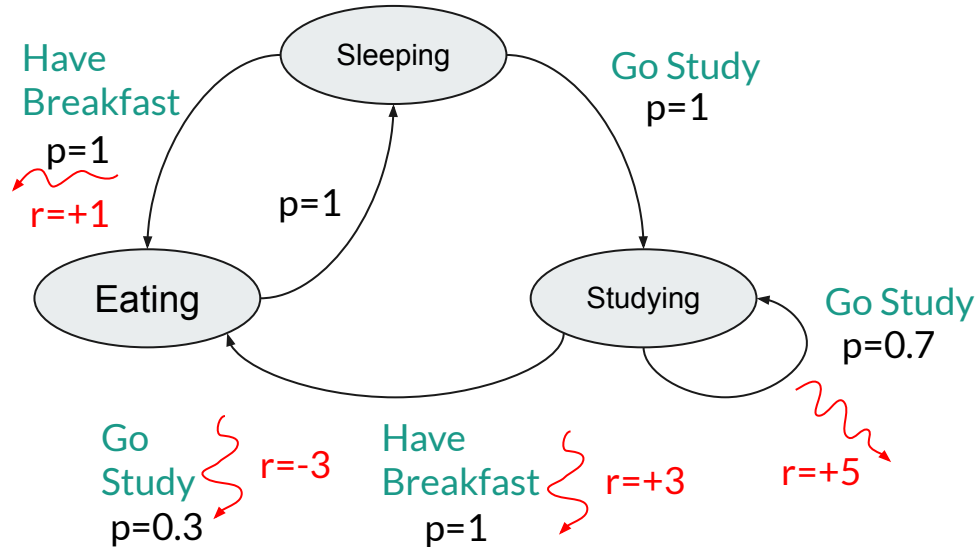
$(S, A, T, R, \gamma)$



# Markov Decision Process

- States (S)
- Actions (A)
- Transition Probabilities (T/P)
- Rewards (R)
- Discount Factor ( $\gamma$ )

$(S, A, T, R, \gamma)$

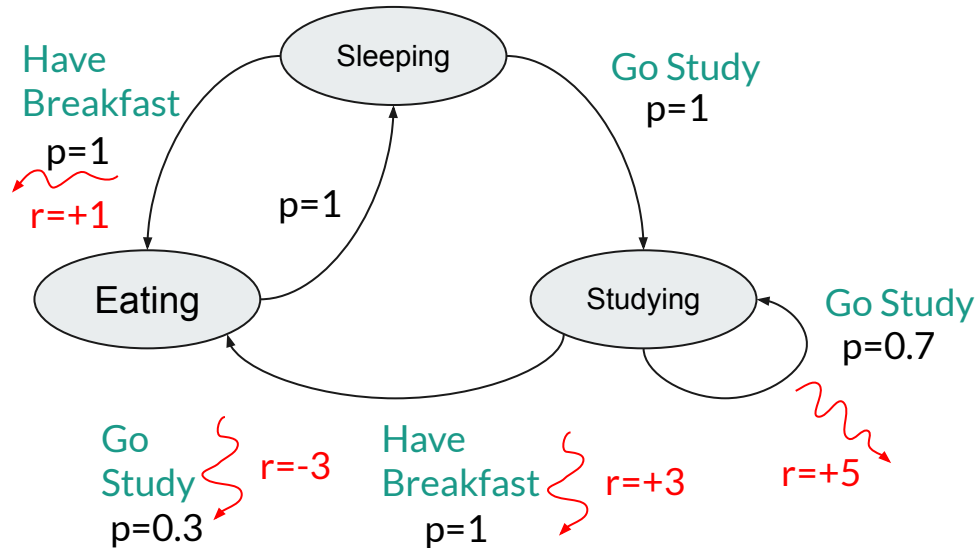


$\pi(s) \rightarrow a$

# Markov Decision Process

- States (S)
- Actions (A)
- Transition Probabilities (T/P)
- Rewards (R)
- Discount Factor ( $\gamma$ )

$(S, A, T, R, \gamma)$




$\pi(s) \rightarrow a$

Deterministic or  
stochastic  
environment?



# Discount Factor & Return

$\tau_1$




$\tau_2$



$$G_t = r_{t+1} + \gamma r_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

## Discount Factor & Return




If  $\gamma=1$ :  $G = 1*0 + 1*3 + 1*0 = 3$



If  $\gamma=1$ :  $G = 1*1 + 1*0 = 1$

$$G_t = r_{t+1} + \gamma r_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

## Discount Factor & Return




If  $\gamma=0$ :       $G =$       ???



If  $\gamma=0$ :       $G =$       ???

$$G_t = r_{t+1} + \gamma r_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

## Discount Factor & Return



If  $\gamma=0$ :  $G = (0^0)*0 + (0^1)*3 + (0^2)*0 = 0$

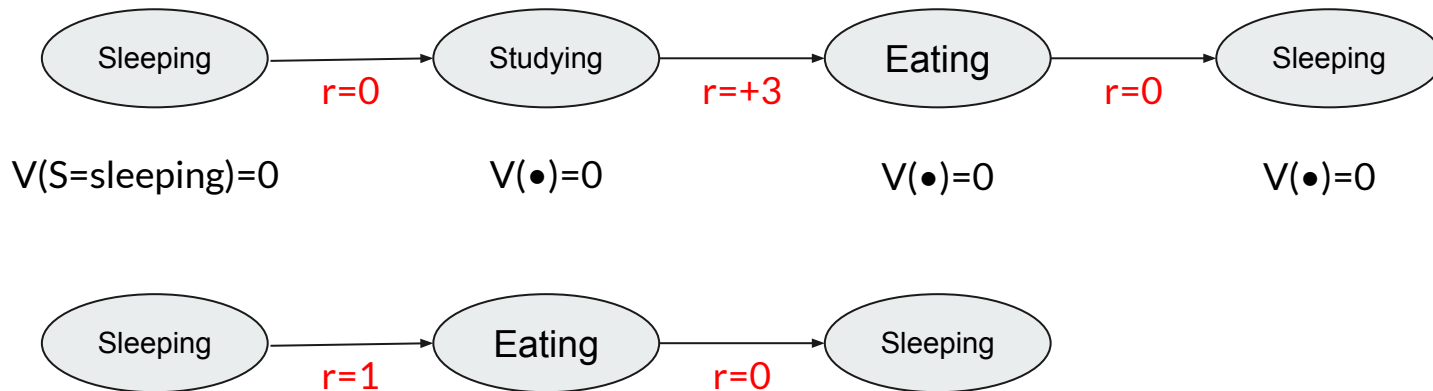


If  $\gamma=0$ :  $G = (0^0)*1 + (0^1)*0 = 1$

$$V^\pi(s) = \mathbb{E}_\pi[G_t | s_t = s] = \mathbb{E}_\pi[r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots | s_t = s]$$

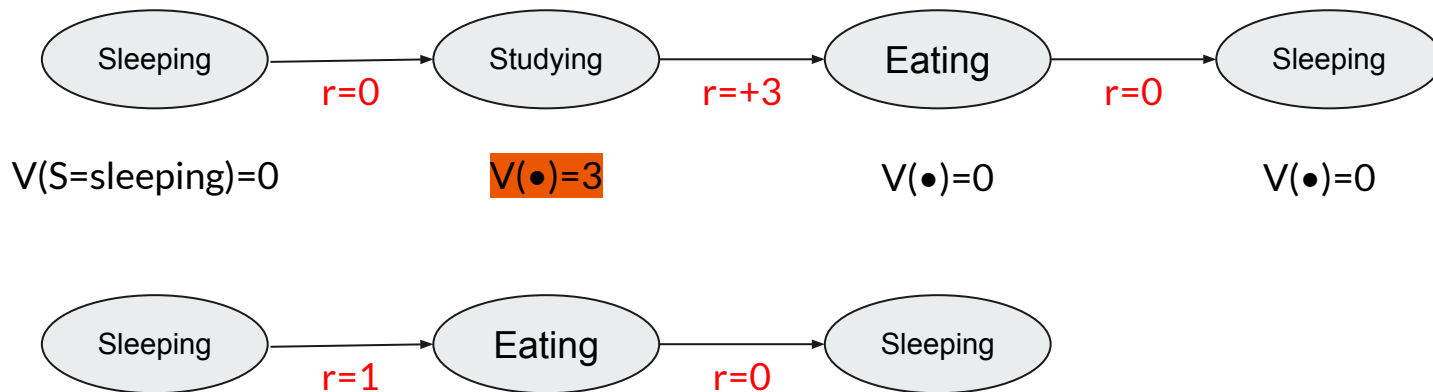
“How good is a state?”

## Value Function



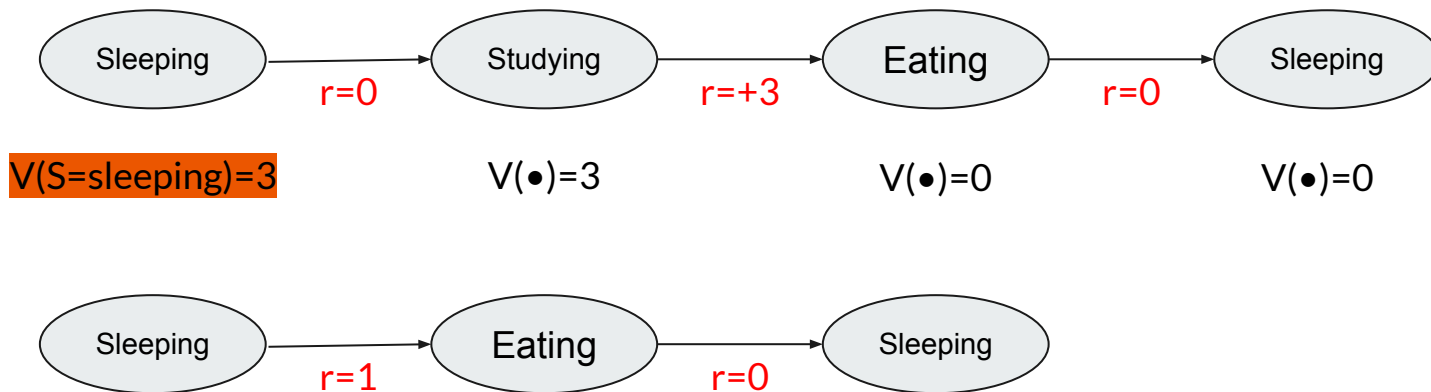
$$V^\pi(s) = \mathbb{E}_\pi[G_t | s_t = s] = \mathbb{E}_\pi[r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots | s_t = s]$$

## Value Function



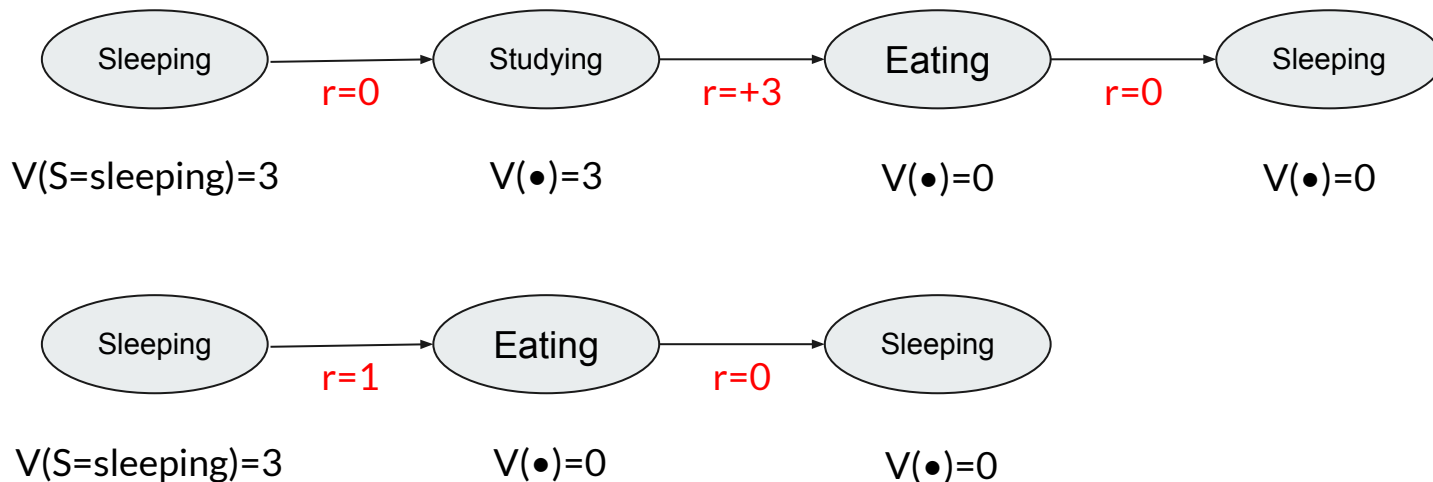
$$V^\pi(s) = \mathbb{E}_\pi[G_t | s_t = s] = \mathbb{E}_\pi[r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots | s_t = s]$$

## Value Function



$$V^\pi(s) = \mathbb{E}_\pi[G_t | s_t = s] = \mathbb{E}_\pi[r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots | s_t = s]$$

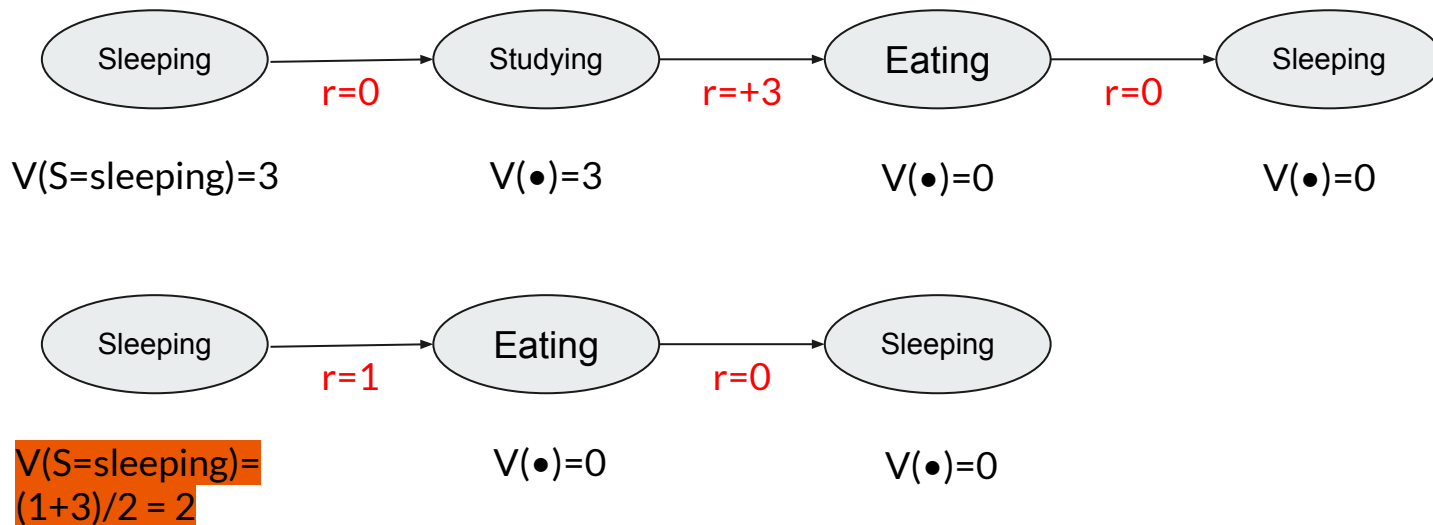
## Value Function





$$V^\pi(s) = \mathbb{E}_\pi[G_t | s_t = s] = \mathbb{E}_\pi[r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots | s_t = s]$$

## Value Function



# Bellman Equation

$$V^{\pi}(s) = R(s, \pi(s)) + \gamma \sum_{s'} T(s, \pi(s), s') V^{\pi}(s')$$

Value of  $s$       Immediate reward       $s'$       Probability of reaching  $s'$  from  $s$  under  $\pi$       Value of  $s'$

---

## Part 2 - Model-based RL



## Model-based RL Idea


1. Build (or have) a model of the environment dynamics: transitions (T), rewards (R)
2. “Solve” the environment: via planning or value estimation
3. ...
4. Profit



## Learning the Model

T & R can be learned by averaging observations from trajectories.

## Input Policy $\pi$



Assume:  $\gamma = 1$

## Observed Episodes (Training)

### Episode 1

B, east, C, -1  
C, east, D, -1  
D, exit, x, +10

### Episode 2

B, east, C, -1  
C, east, D, -1  
D, exit, x, +10

### Episode 3

E, north, C, -1  
C, east, D, -1  
D, exit, x, +10

### Episode 4

E, north, C, -1  
C, east, A, -1  
A, exit, x, -10

## Learned Model

$$\hat{T}(s, a, s')$$

T(B, east, C) =   
T(C, east, D) =   
T(C, east, A) =   
...

$$\hat{R}(s, a, s')$$

R(B, east, C) =   
R(C, east, D) =   
R(D, exit, x) =   
...




# Policy Evaluation

(Goal: get value function everywhere)

1. Initialize VF everywhere at 0
2. Iterate over all states; update their value with the reward from all reachable states + their current value function
3. GOTO 2

# Policy Evaluation




Assume  $T(\bullet) = 1$ ,  $\gamma = 1$ ,  
uniform random policy

	0	
0	0	0
	0	

Iteration 0



# Policy Evaluation




Assume  $T(\bullet) = 1$ ,  $\gamma = 1$ ,  
uniform random policy

	-10	
-1	-1	+10
	-1	

Iteration 1

# Policy Evaluation




Assume  $T(\bullet) = 1$ ,  $\gamma = 1$ ,  
uniform random policy

	-10	
-1+(-1)	$\frac{1}{2}$ $(-1-10)+\frac{1}{2}$ $(-1+10) =$ -1	+10
	-1+(-1)	

Iteration 2

# Policy Evaluation




Assume  $T(\bullet) = 1$ ,  $\gamma = 1$ ,  
uniform random policy

	-10	
-2	-1	+10
	-2	

Iteration 2

## And another one



- Undiscounted episodic MDP ( $\gamma = 1$ )
- Nonterminal states  $1, \dots, 14$
- One terminal state (shown twice as shaded squares)
- Actions leading out of the grid leave state unchanged
- Reward is  $-1$  until the terminal state is reached
- Agent follows uniform random policy

$$\pi(n|\cdot) = \pi(e|\cdot) = \pi(s|\cdot) = \pi(w|\cdot) = 0.25$$

# $v_k$ for the Random Policy

# Greedy Policy w.r.t. $v_k$

$k = 0$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

	↕↕↕	↕↕↕	↕↕↕
↕↕↕	↕↕↕	↕↕↕	↕↕↕
↕↕↕	↕↕↕	↕↕↕	↕↕↕
↕↕↕	↕↕↕	↕↕↕	

random  
policy

$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

	←	↕↕↕	↕↕↕
↑	↕↕↕	↕↕↕	↕↕↕
↕↕↕	↕↕↕	↕↕↕	↓
↕↕↕	↕↕↕	→	

$$\begin{aligned} & \frac{1}{4} * (-1 + 0) + \\ & \frac{1}{4} * (-1 - 1) + \\ & \frac{1}{4} * (-1 - 1) + \\ & \frac{1}{4} * (-1 - 1) \end{aligned}$$

Chance of going there under  
current policy \*  $k = 2$   
(Reward of transition +  
 $V(s')$ )

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

	←	←	↕↕↕
↑	↖	↕↕↕	↓
↑	↕↕↕	↗	↓
↕↕↕	→	→	

$v_k$  for the  
Random Policy

Greedy Policy  
w.r.t.  $v_k$

$k = 0$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

	↕↕↕	↕↕↕	↕↕↕
↕↕↕	↕↕↕	↕↕↕	↕↕↕
↕↕↕	↕↕↕	↕↕↕	↕↕↕
↕↕↕	↕↕↕	↕↕↕	

← random  
policy

$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

	←	↕↕↕	↕↕↕
↑	↕↕↕	↕↕↕	↕↕↕
↕↕↕	↕↕↕	↕↕↕	↓
↕↕↕	↕↕↕	→	

$k = 2$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

	←	←	↕↕↕
↑	↖	↕↕↕	↓
↑	↕↕↕	↘	↓
↕↕↕	→	→	

$k = 3$

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

	←	←	↙
↑	↖	↘	↓
↑	↖	↘	↓
↖	→	→	

$k = 10$

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0


	←	←	↙
↑	↖	↘	↓
↑	↖	↘	↓
↖	→	→	

$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0


	←	←	↙
↑	↖	↘	↓
↑	↖	↘	↓
↖	→	→	

optimal policy



**Policy evaluation** Estimate  $v_\pi$   
 Iterative policy evaluation

**Policy improvement** Generate  $\pi' \geq \pi$   
 Greedy policy improvement







## Problems:

- Costly (synchronous update of all states + every state vs. every accessible state)
- Need transition function

→ how about asynchronous updates as we go? (MC/TD(0))

---

## Part 2 - Model-free RL



# Monte-Carlo Policy Evaluation

Idea:

- When an episode is over, store actual mean return ( $G$ ) for each state
- Update value function to approximate this  $G$  for each state

Learning Rate

$$V(S_t) \leftarrow V(S_t) + \boxed{\alpha} (G_t - V(S_t))$$




# Incremental Updates

New value = old value + learning rate \* (measurement - old value)

If measurement == old value, then no change,  
otherwise small increase/decrease

# MCPE




Assume  $\gamma = 1, \alpha = 0.1$

	0	
0	0	0
	0	

Iteration 0

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$

## MCPE



Assume  $\gamma = 1, \alpha = 0.1$

	0	
0	0	0
	0	

Iteration 0


Trajectory 1

B → C, -1  
 C → D, -1  
 D → x, +10

G(B): +8  
 G(C): +9  
 G(D): +10

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$

## MCPE



Assume  $\gamma = 1, \alpha = 0.1$

	0	
$0 + 0.1 * 8 = .8$	$0 + 0.1 * 9 = .9$	$0 + 0.1 * 10 = 1$
	0	

Iteration 1


Trajectory 1

B → C, -1  
C → D, -1  
D → x, +10

G(B): +8  
G(C): +9  
G(D): +10

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$

## MCPE



Assume  $\gamma = 1, \alpha = 0.1$

	0	
.8	.9	1
	0	

Iteration 1

Trajectory 1    Trajectory 2

B→C, -1  
C→D, -1  
D→x, +10

B→C, -1  
C→A, -1  
A→x, -10


G(B): +8  
G(C): +9  
G(D): +10

G(B): -12  
G(C): -11  
G(A): -10



$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$

## MCPE



Assume  $\gamma = 1, \alpha = 0.1$

	$0 + .1 * -10$ $= -.1$	
$.8 + .1 * (-12 - .8)$ $= -0.48$	$.9 + .1 * (-11 - .9)$ $= -0.29$	1
	0	

Iteration 2

Trajectory 1    Trajectory 2

B → C, -1  
C → D, -1  
D → x, +10

B → C, -1  
C → A, -1  
A → x, -10

G(B): +8  
G(C): +9  
G(D): +10

G(B): -12  
G(C): -11  
G(A): -10



# Temporal Difference Learning

Idea:

- Same thing but we don't wait for the episode's end
- Use single step (reward+V(s')) to update  $\rightarrow$  single step = TD(0)

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$



$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

Source: David Silver reinforcement learning lecture series,  
<http://www0.cs.ucl.ac.uk/staff/d.silver/web/Teaching.html>,

---

# Part 3 - RL for Control



## Summary so far

- Learned about MDP/ $(S, A, T, R, \gamma)$
- Policy Eval learns value func. (given  $T, R, \pi$ )
- Monte-Carlo Policy Eval learns value func (given  $\pi$ )
- Temporal Difference Learning learns value func (given  $\pi$ )
- Can use greedy  $\pi$  if we have  $T$ , but what if we don't?

→ Q function to the rescue



# Q Learning

Similar to value function, but also taking actions into consideration:

## Definition

The *action-value function*  $q_{\pi}(s, a)$  is the expected return starting from state  $s$ , taking action  $a$ , and then following policy  $\pi$

$$q_{\pi}(s, a) = \mathbb{E}_{\pi} [G_t \mid S_t = s, A_t = a]$$



## Q Learning Policy

$$\pi'(s) = \operatorname{argmax}_{a \in \mathcal{A}} q_{\pi}(s, a)$$

At each state, check the Q value of all the actions; Pick action with highest Q



## How to learn Q?

- Use Monte-Carlo algorithm (“Monte-Carlo Q Learning”)
- Use Temporal Difference algorithm (“Sarsa”) - same as TD(0) but with Q function instead of value function



## Monte-Carlo Q Learning

- For each state  $S_t$  and action  $A_t$  in the episode,


$$N(S_t, A_t) \leftarrow N(S_t, A_t) + 1$$

Simple  
counter  
(start at 0)

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{1}{N(S_t, A_t)} (G_t - Q(S_t, A_t))$$



# Sarsa - State-action-reward-state-action



$$Q(S, A) \leftarrow Q(S, A) + \alpha (R + \gamma Q(S', A') - Q(S, A))$$



**BUT: exploration**



"Behind one door is tenure - behind the other is flipping burgers at McDonald's."

- There are two doors in front of you.
- You open the left door and get reward 0  
 $V(\text{left}) = 0$
- You open the right door and get reward +1  
 $V(\text{right}) = +1$
- You open the right door and get reward +3  
 $V(\text{right}) = +2$
- You open the right door and get reward +2  
 $V(\text{right}) = +2$
- $\vdots$
- Are you sure you've chosen the best door?



## $\epsilon$ -Greedy Exploration

- With probability  $1 - \epsilon$  choose the greedy action
- With probability  $\epsilon$  choose an action at random

You can decrease  $\epsilon$  slowly as you go



## How to Neural-network this?

- Q function: a neural network, in: observation, action, out: scalar float
- $\pi$ : testing each action + picking highest Q value
- “Actor-Critic” = Q network + policy network

Continuous actions:

- $Q(s,a) = Q(s, \pi(s)) \rightarrow$  backprop through Q into both nets

---

# Part 4 - Practical RL



# OpenAI gym

- De-facto standard RL environment(s)
- Contains variety of tasks (text adventures, Atari games, robotic tasks...)

→ show [https://gym.openai.com/envs/#classic\\_control](https://gym.openai.com/envs/#classic_control)

- Only tasks, no learning algorithms
- Homogenous API

```
1  import gym
2
3  env = gym.make("Pendulum-v0")
4
5  obs = env.reset()
6
7  # optional
8  env.render()
9
10 done = False
11
12 while not done:
13     # in practice the action comes from your policy
14     action = env.action_space.sample()
15
16     obs, rew, done, misc = env.step(action)
17
18     # optional
19     env.render()
20
21
```



```
obs, rew, done, misc = env.step(action)
```

Observation as  
List, Tuple, Numpy array

Ex: img,  
`np.array((128,128,3),  
dtype=np.uint8)`

Ex: robot joints + velocities  
`list(0.4, 1.0, -0.3, 0.0)`

Scalar float

Ex: +10  
Ex: -0.001

Bool

Ex: True  
Ex: False

Dictionary

Ex: {}

Ex:  
{“success”: True,  
“steps”: 420}

Ex:  
{“reward\_pos”: 69,  
“reward\_vel”: 1,  
“reward\_rules”: -10.4}

```
1 import gym
2
3 env = gym.make("Pendulum-v0")
4
5 policy = Policy() # <-- not part of Gym
6 replay_buf = ReplayBuffer() # <-- also not part of Gym
7
8 while True:
9
10     obs = env.reset()
11     done = False
12
13     while not done:
14
15         action = policy.select_action(obs)
16
17         new_obs, rew, done, misc = env.step(action)
18
19         replay_buf.add((obs, action, new_obs, rew, done))
20
21     replay_batch = replay_buf.sample()
22     policy.train(replay_batch)
23
```

Find  
the  
error!



# Best Practices

- Normalize observations & actions to be in  $[-1,1]$  or  $[0,1]$ :

`np.array(-100, 5, 30) → np.array(-1, 0.05, .3)`

(normalize by max/range or by mean/std)



# Best Practices

- Normalize observations & actions to be in  $[-1,1]$  or  $[0,1]$ :

`np.array(-100, 5, 30) → np.array(-1, 0.05, .3)`

(normalize by max/range or by mean/std)

- Limit/scale rewards

-100 on failure, +1 on success, -0.00001 per step → -5, +1, -0.01



# Best Practices

- Normalize observations & actions to be in  $[-1,1]$  or  $[0,1]$ :

`np.array(-100, 5, 30) → np.array(-1, 0.05, .3)`

(normalize by max/range or by mean/std)

- Limit/scale rewards

-100 on failure, +1 on success, -0.00001 per step → -5, +1, -0.01

- Make sure environment is markovian



# Best Practices

- Normalize observations & actions to be in  $[-1,1]$  or  $[0,1]$ :

`np.array(-100, 5, 30) → np.array(-1, 0.05, .3)`

(normalize by max/range or by mean/std)


- Limit/scale rewards

-100 on failure, +1 on success, -0.00001 per step → -5, +1, -0.01

- Make sure environment is markovian

And  
LOTS OF SEEDS

# Reproducibility





# Markov property

Example:

**Flying a helicopter**

Observation:

(position\_xyz, velocity\_xyz)





# Markov property

Example:

**Flying a helicopter**

Observation:

`(position_xyz, velocity_xyz)`

Better:

`(position_xyz, velocity_xyz,  
rotation_quat, goal_xyz)`



# Markov property

Example:

## Flying a helicopter

Observation:

`(position_xyz, velocity_xyz)`

Better:

`(position_xyz, velocity_xyz,  
rotation_quat, goal_xyz)`

Example:

## Driving a car

Observation:

`Image, (256, 256, 3)`



# Markov property

Example:

## Flying a helicopter

Observation:

`(position_xyz, velocity_xyz)`

Better:

`(position_xyz, velocity_xyz,  
rotation_quat, goal_xyz)`

Example:

## Driving a car

Observation:

`Image, (256, 256, 3)`

Better:

`Image stack (last 4 images) + depth images  
(4, 256, 256, 3) + (4, 256, 256, 1)`

Implement stuff      Tweak until it works  
(80 hours)              (40 hours)



Here's how long each stage *actually* took.

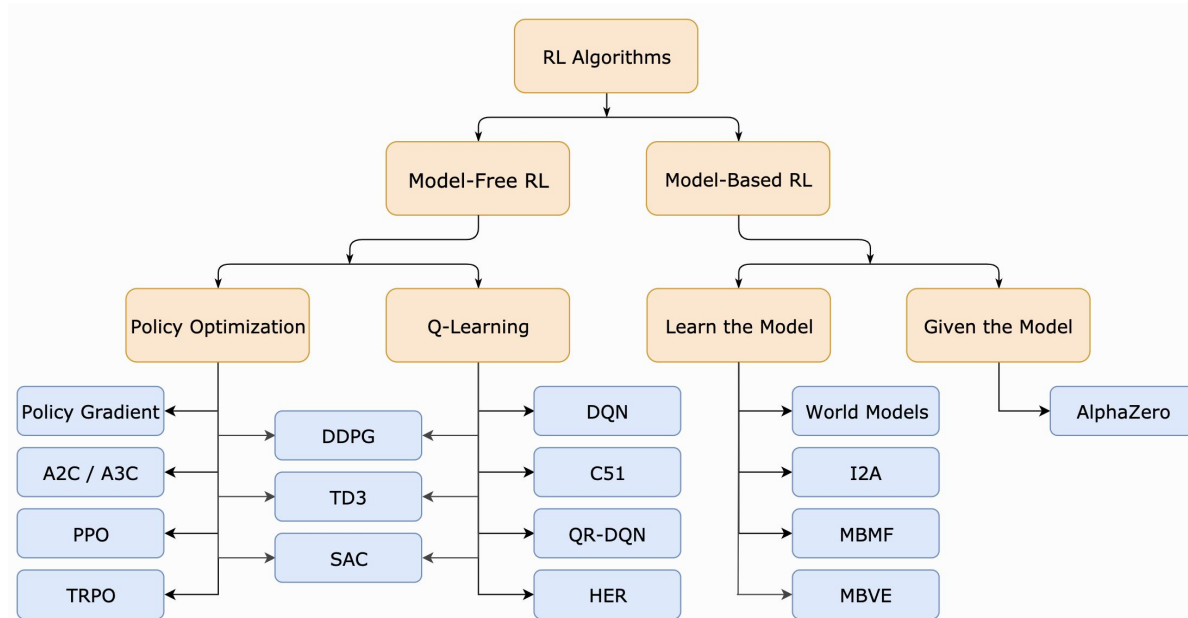
Implement stuff              Get it working              Get reliable tests working,  
with a toy environment              clean up code  
(30 hours)              (110 hours)              (60 hours)



Get it working  
with Pong/Enduro  
(10/10 hours)



# State-of-the-Art DRL algos



Source: OpenAI Spinning Up,

[https://spinningup.openai.com/en/latest/spinningup/rl\\_intro2.html](https://spinningup.openai.com/en/latest/spinningup/rl_intro2.html)




## How to pick?


9/10 times: PPO works (discrete or continuous actions)

+ few HP adjustments (episodes, stacked frames, hidden rep size)

Otherwise, try SAC (slow but stable) or TD3 (fast, easy, sensitive to seed)



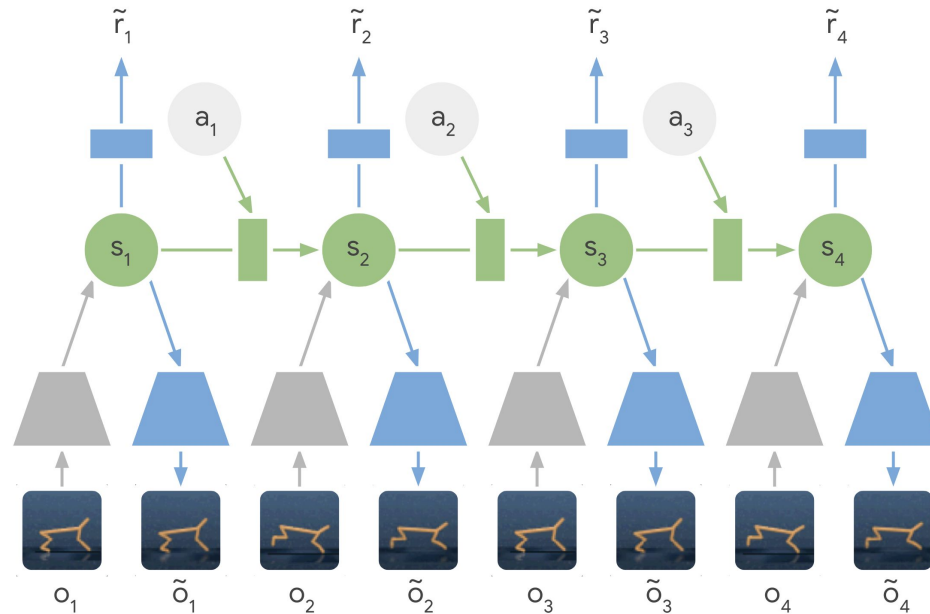
X axis: steps in the environment  
(every time the “env.step()” function is called)



X axis: compute time in seconds



## Bonus: Modern Planning via PlaNet





## Bonus: Imitation Learning

Naive: behavior cloning

- Can be used as initialization for RL policy
- But overfits to training data

Better: DAgger / Deeply AggreVaTeD, see

[http://videolectures.net/DLRLsummerschool2018\\_daume\\_imitation\\_learning/](http://videolectures.net/DLRLsummerschool2018_daume_imitation_learning/)

# Thanks, questions?

---